

Interview Insights Report

Executive Summary

Across the fourteen interviews conducted, a consistent picture emerges of employees who are willing to use AI but who make split-second, practical calculations about whether it is worth their time. Adoption tends to stick when AI removes a clear, high-friction task, such as pulling together a status update across email, chat, and notes, or turning a raw data set into a working script. Adoption breaks down when the tool requires heavy babysitting, when it produces something that misses the point entirely, or when the interface itself gets in the way of a visual or fast-moving workflow. Trust is fragile and asymmetric: a single bad experience with a high-stakes task can end usage for that category permanently, while good experiences build confidence slowly and unevenly across teams. Underneath the individual behavior sits an organizational layer that is still finding its footing, with training that can feel scattered at times, incentives that vary by department, and pockets of informal, peer-led knowledge sharing that seem to be doing more for real adoption than top-down programs thus far. The insights below are grouped into behavioral barriers, prompt engineering and technical patterns, and ecosystem and training dynamics, with supporting detail and interviewee attribution under each theme.

1 - Behavioral Barriers

ROI and Time Investment

Employees are running a constant, informal cost-benefit analysis on whether prompting is faster than just doing the work themselves.

- Tasks that take less than 10 to 20 minutes manually get skipped entirely if the prompting process requires multiple rounds of correction, since the back and forth erases any time saved (Interviewee 2).
- Spending 4 hours setting up or refining a prompt workflow to save just 1 hour of manual labor is viewed as a bad trade, especially on a packed schedule (Interviewee 9).
- Attempting to automate a project manager workflow backfired for one interviewee, who ended up spending more time troubleshooting prompts than the analysis would have taken manually (Interviewee 8).

- High-frequency administrative work, like manual 217-row mandate setups that take 30 to 60 minutes each, creates real appetite for automation, but the learning curve before seeing a payoff is steep enough to discourage people from starting (Interviewee 9).
- When a prompt fails or feels like too much effort, the instinct is to abandon it and go back to the manual method rather than keep troubleshooting (Interviewee 13).
- Constantly having to self-correct an AI's output is enough reason to just do the task by hand instead (Interviewee 4).
- Even for someone who uses AI roughly half the day, writing tasks still split close to evenly between AI use and manual work, since refining and fact-checking the output can eat up nearly as much time as generating it (Interviewee 14).

Trust Degradation from Intent Distortion and Performance Inaccuracies

Automated changes or errors that miss the point of the task do lasting damage to confidence in the tool.

- A built-in "polish" feature that missed the user's actual arguments led to it being avoided altogether, along with a warning that AI is set up to pull people into distracting rabbit holes instead of keeping them on task (Interviewee 6).
- Unsolicited clarifying questions from an LLM can feel intrusive. Several users would rather get a direct answer without the extra back and forth (Interviewee 1).
- Trust breaks down when a tool starts adding unrequested details or generating what one interviewee called nonsense that has nothing to do with the specific task, leading to avoidance for specialized work (Interviewee 5).
- A single high-stakes mistake can trigger long-term avoidance of a tool in critical areas, since the person starts questioning whether it can be trusted again at all (Interviewee 1, Interviewee 6).
- Hallucination and flat-out misunderstanding of a request is a real concern, and it is one of the main reasons a user will abandon an AI-assisted task partway through (Interviewee 4).
- A lingering worry that AI might quietly overlook details in math, data comparisons, or legal language means some tasks are done manually from the start, since double-checking the AI would take just as long (Interviewee 13).

Visual and Interface Friction

Text-first tools create friction for design-heavy or visually oriented work.

- Visual learners find text-based consoles intimidating and get frustrated when they cannot directly drag, drop, or tweak AI-generated assets in real time (Interviewee 9).

- Current LLM tools produce weak slide deck design, forcing users to manually redo layout and formatting before it is fit for senior leadership (Interviewee 5, Interviewee 6).
- The manual steps involved in using the internal LLM, finding a file, dropping it in, and switching between browser tabs, feel tedious when trying to move quickly, which is part of why quick lookups often happen on a regular browser instead (Interviewee 12).
- Needing to chain outputs from one AI tool into another adds enough friction that one interviewee is building a personal tool just to solve that handoff problem (Interviewee 14).

Voluntary and Discretionary Boundary Restrictions

Employees are setting their own limits to manage data security and accuracy risk, often without being told to.

- Users are careful and strict about inputting sensitive client data, such as Excel sheets with user data, user IDs, or account numbers, into the internal LLM Suite (Interviewee 2, Interviewee 9).
- Some manually sanitize data before feeding it into an AI tool, scrubbing out names and personal details as a precaution (Interviewee 9).
- AI is deliberately kept out of calculations and other deterministic math tasks because of a general sense that LLMs are not reliable with numbers (Interviewee 5, Interviewee 6).
- Sensitive information like account numbers is intentionally kept out of prompts, and new chats are frequently started specifically to avoid carrying proprietary context over from one conversation to the next (Interviewee 13).
- Some users describe themselves as naturally conservative and protective of what they input, independent of any specific bad experience, and note that formal compliance training on AI use is currently ramping up (Interviewee 12).

Confidence Among Lower-Friction Users

Not every interviewee described caution as the default. A subset of employees, often in more technical or build-facing roles, reported little to no hesitation about security or compliance, which suggests risk tolerance is not uniform.

- One interviewee reported no compliance or security concerns at all, attributing it to working inside a highly regulated institution with guardrails already built in, which frees them to use AI for essentially everything (Interviewee 4).
- Another echoed the same reasoning almost verbatim, citing the regulated environment as the reason they do not hold back on AI use for their work (Interviewee 14).

- Trust in the firm's internal systems was enough for one interviewee to feel comfortable with AI generally, even while still double-checking outputs before sending them onward (Interviewee 10).

2 - Prompt Engineering and Technical Insights

Lack of Internal Context

Generic prompting struggles because the AI does not natively know the company's internal shorthand or specific data.

- LLMs frequently hallucinate localized details, including internal terminology, speaker identity in meeting transcripts, and specific project locations (Interviewee 1).
- Users get worn down by the upfront effort of re-teaching company-specific context and siloed data to internal bots at the start of every new session before getting anything usable back (Interviewee 9).
- When enough explicit context is included directly in the prompt, the LLM can still land the right answer even from a garbled or misheard starting point, successfully working out unfamiliar terms and acronyms from a live call based on context alone (Interviewee 12).
- The most-wanted future capability by several interviewees is a persistent, personal knowledge base that understands everything going on across a person's apps and history, rather than an assistant that starts from zero each time (Interviewee 14).

Leveraging Intermediary Code Generation

Prompting the AI to build a tool or script tends to outperform asking it to process raw data directly.

- Complex data reconciliation across massive data sets went from days to about one hour by prompting the LLM to generate Python code or a macro instead of asking it to do the analysis itself (Interviewee 8).
- Prompting an LLM to generate advanced, recursive data validation formulas lets a non-technical employee skip traditional search entirely and execute a technical task they would not have known how to do by hand (Interviewee 3).
- Prior coding experience is not a prerequisite for high-impact automation. One user with only an introductory Python class produced a working script using AI (Interviewee 8).
- Existing code-based automation still has real limits for non-technical users: an agenda-building script pulls from a shared Excel tracker automatically, but adding new

team members to it means going directly into the code to update it manually, which is a barrier for anyone without a coding background (Interviewee 10).

- Screenshotting a workflow and prompting an AI coding tool for specific changes was used effectively to speed through prototyping and platform handoff work (Interviewee 14).

Cross-Platform Data Aggregation

The clearest, highest-satisfaction value shows up in prompts that synthesize information across different communication tools.

- Prompts that pull data across emails, chats, and shared notes into one unified status summary consistently satisfy users and produce clean, multi-layered executive overviews (Interviewee 6).
- The top capability employees want is an AI that can see across every core application at once, email, spreadsheets, calendar, and presentations, to close the gap between disconnected tools (Interviewee 1).
- Meaningful time is saved when prompts consolidate separate meeting note documents, automatically surfacing discrepancies and key differences between them (Interviewee 5).
- A meeting-recording AI that builds out minutes automatically from a call is seen as a genuinely universal use case, one that nearly any team could adopt in the same way (Interviewee 10).
- There is a clear desire for AI that bridges disconnected internal software entirely, such as syncing a workflow case management tool with email so information does not have to be manually cross-checked between them (Interviewee 13).

Low-Friction Verification Mechanics

Trust depends less on understanding the AI's internal logic and more on having a fast way to spot-check its output.

- Full transparency into the AI's backend reasoning is not what users want. A fast, independent way to sample and verify results, like spot-checking around 200 trades out of thousands and isolating the mismatches, is what actually builds confidence (Interviewee 8).
- Line-by-line textual verification can be skipped by cross-referencing AI output against an existing visual system, like a color-coded weekly tracker, to quickly confirm accuracy (Interviewee 5).
- One interviewee's version of verification was to rerun their own personal notes through the LLM separately and compare that output against the AI-generated meeting minutes before sending a final version to the wider group (Interviewee 10).

- Verification does not always need to be heavy. Sometimes it is simply reading the output once to confirm it matches the intended message before sending it along (Interviewee 13).

3 - Ecosystem and Training Dynamics

Mandated Accountability

Structural incentives appear to be driving daily engagement faster than voluntary effort alone.

- Tying AI usage metrics directly to performance evaluations helped to remove organizational pushback and forced daily interaction within more structured departments (Interviewee 8).
- Usage is tracked at the leadership level in some teams even without a mandatory quota attached, which still seems to nudge adoption without formal pressure (Interviewee 13).
- For at least one interviewee, the push to use AI comes across less as encouragement and more as a competitive necessity, framed as something you cannot afford to fall behind on (Interviewee 14).

Training Fatigue

Too many training options without role-specific direction leaves people stuck rather than supported.

- Training invitations create a paradox of choice, and users end up wishing for direction specific to their role instead of broad, generic tutorials (Interviewee 9).
- Firmwide voluntary pushes can feel like checking a box, producing little real skill growth compared to what people pick up through their own trial and error and specific prompt engineering instruction (Interviewee 3).
- Fragmented tools across different lines of business create siloed workflows and a strong desire for one consistent streamlined aggregation (Interviewee 2).

Grassroots Knowledge Networks

Informal, localized sharing appears to scale tactical prompting knowledge more effectively than formal programs.

- Casual, global team office hours built specifically to share shortcuts and failures help stop team members from constantly “reinventing the wheel” across regional divides (Interviewee 6).

- A dedicated tech champion within a team, sending targeted newsletters and custom agent instructions, keeps larger groups aligned on what tools are relevant to them (Interviewee 2).
- A designated tech expert coordinating globally, for example with the EMEA region, accelerates localized learning by sharing concrete workflow efficiencies and assignments (Interviewee 5).
- Regular twice-weekly team meetings function as an ongoing forum where anyone can share a new discovery or implementation, without needing a formal program behind it (Interviewee 10).
- Peer-to-peer sharing during regular meetings, paired with self-paced instructional videos pushed down from corporate, forms the backbone of how one team learns, with almost no mandatory structure involved (Interviewee 13).
- Biweekly shareouts and simple, informal conversation between teammates are how one product team keeps everyone current, more so than any structured training (Interviewee 14).